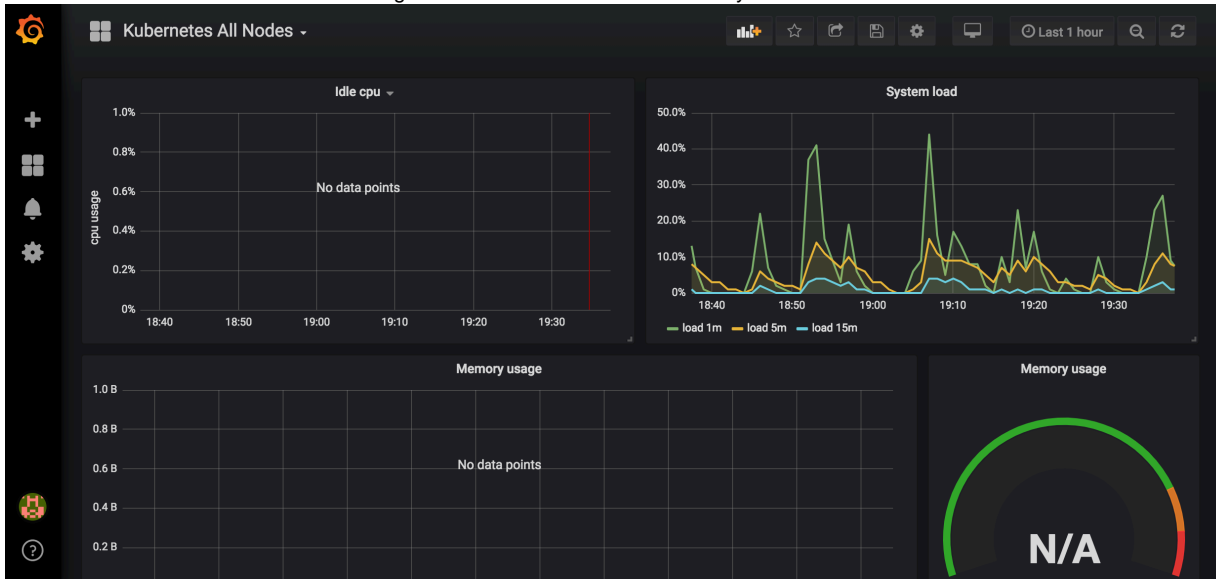


Deploying Prometheus and Grafana in a Gospel EKS cluster

- 1. Pre-requisites
- 2. Deploying Prometheus
- 3. Deploying Grafana
- 4. Creating Grafana dashboards
 - Source for these instructions - AWS EKS workshop

We will be using Prometheus and Grafana to monitor Gospel Kubernetes clusters, the audience for these dashboards being system/database administrators that are interested in having an overview of the health of their system.



1. Pre-requisites

Helm - you can set it up with these steps

https://eksworkshop.com/beginner/060_helm/helm_intro/install/

An EKS cluster running Gospel in its own separate

namespace

2. Deploying Prometheus

Prometheus is used in order to gather metrics relevant for the Kubernetes cluster which will ultimately be used by Grafana to graph data.

```
kubectl create namespace prometheus

helm install stable/prometheus \
  --name prometheus \
  --namespace prometheus \
  --set alertmanager.persistentVolume.storageClass="gp2" \
  --set server.persistentVolume.storageClass="gp2"
```

Make note of the prometheus endpoint in helm response (you will need this later). It should look similar to below:

The Prometheus server can be accessed via port 80 on the following DNS name from within your cluster:
prometheus-server.prometheus.svc.cluster.local

Check if Prometheus components deployed as expected

```
kubectl get all -n prometheus
```

You should see response similar to below. They should all be Ready and Available

NAME	READY	STATUS
pod/prometheus-alertmanager-77cfd85db-s9p48	2/2	Running
0 1m		
pod/prometheus-kube-state-metrics-74d5c694c7-vqtjd	1/1	Running
0 1m		
pod/prometheus-node-exporter-6dhpw	1/1	Running
0 1m		
pod/prometheus-node-exporter-nrfkn	1/1	Running
0 1m		
pod/prometheus-node-exporter-rtrm8	1/1	Running
0 1m		
pod/prometheus-pushgateway-d5fdc4f5b-dbmrg	1/1	Running
0 1m		
pod/prometheus-server-6d665b876-dsmh9	2/2	Running
0 1m		

NAME	TYPE	CLUSTER-IP
EXTERNAL-IP PORT(S) AGE		
service/prometheus-alertmanager	ClusterIP	10.100.89.154
<none> 80/TCP 1m		
service/prometheus-kube-state-metrics	ClusterIP	None
<none> 80/TCP 1m		
service/prometheus-node-exporter	ClusterIP	None
<none> 9100/TCP 1m		
service/prometheus-pushgateway	ClusterIP	10.100.136.143
<none> 9091/TCP 1m		
service/prometheus-server	ClusterIP	10.100.151.245
<none> 80/TCP 1m		

NAME	DESIRED	CURRENT	READY
UP-TO-DATE AVAILABLE NODE SELECTOR AGE			
daemonset.apps/prometheus-node-exporter	3	3	3
3 3 <none> 1m			

NAME	DESIRED	CURRENT	READY
UP-TO-DATE AVAILABLE AGE			
deployment.apps/prometheus-alertmanager	1	1	1
1 1m			
deployment.apps/prometheus-kube-state-metrics	1	1	1
1 1m			
deployment.apps/prometheus-pushgateway	1	1	1
1 1m			

deployment.apps/prometheus-server			1	1	1
1	1m				
NAME			DESIRED		
CURRENT	READY	AGE			
replicaset.apps/prometheus-alertmanager-77cfd85db			1		1
1	1m				
replicaset.apps/prometheus-kube-state-metrics-74d5c694c7			1		1
1	1m				
replicaset.apps/prometheus-pushgateway-d5fdc4f5b			1		1
1	1m				

replicaset.apps/prometheus-server-6d665b876	1	1
1	1m	

3. Deploying Grafana

As with Prometheus, we are setting the storage class to gp2, admin password, configuring the datasource to point to Prometheus and creating an external load balancer for the service.

```
kubectl create namespace grafana
helm install stable/grafana \
  --name grafana \
  --namespace grafana \
  --set persistence.storageClassName="gp2" \
  --set adminPassword='EKS!sAWSome' \
  --set image.tag=6.6.1-ubuntu \
  --set datasources."datasources\.yaml".apiVersion=1 \
  --set datasources."datasources\.yaml".datasources[0].name=Prometheus
\
  --set datasources."datasources\.yaml".datasources[0].type=prometheus
\
  --set
datasources."datasources\.yaml".datasources[0].url=http://prometheus-ser
ver.prometheus.svc.cluster.local \
  --set datasources."datasources\.yaml".datasources[0].access=proxy \
  --set datasources."datasources\.yaml".datasources[0].isDefault=true
\
  --set service.type=LoadBalancer
```

Run the following command to check if Grafana is deployed properly:

```
kubectl get all -n grafana
```

You should see similar results. They should all be Ready and Available

NAME	READY	STATUS	RESTARTS	AGE
pod/grafana-b9697f8b5-t9w4j	1/1	Running	0	2m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
PORT(S)	AGE		
service/grafana	LoadBalancer	10.100.49.172	
abe57f85de73111e899cf0289f6dc4a4-1343235144.us-west-2.elb.amazonaws.com			
80:31570/TCP	3m		

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.apps/grafana	1	1	1	1
2m				

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/grafana-b9697f8b5	1	1	1	2m

You can get Grafana ELB URL using this command. Copy & Paste the value into browser to access Grafana web UI.

```
export ELB=$(kubectl get svc -n grafana grafana -o
jsonpath='{.status.loadBalancer.ingress[0].hostname}')

echo "http://$ELB"
```

When logging in, use the username **admin** and get the password hash by running the following:

```
kubectl get secret --namespace grafana grafana -o
jsonpath="{.data.admin-password}" | base64 --decode ; echo
```

4. Creating Grafana dashboards

Log into the Grafana dashboard using credentials supplied during configuration (available above).

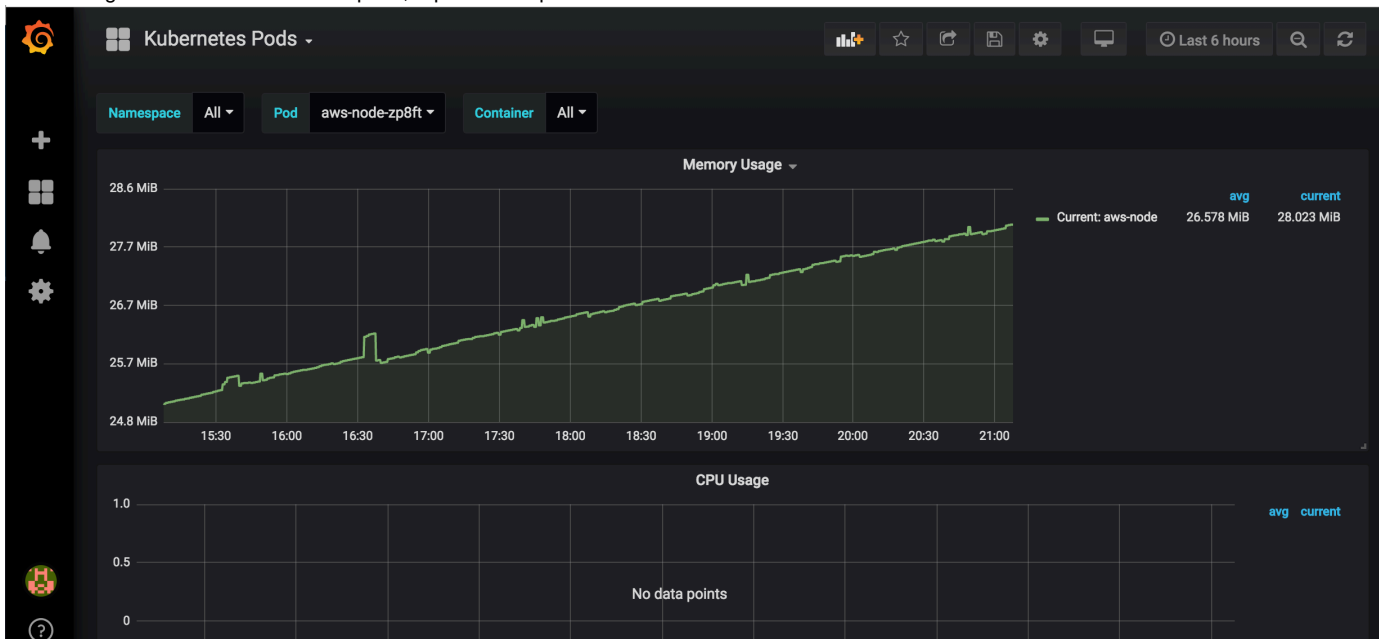
You will notice that **'Install Grafana'** & **'create your first data source'** are already completed. We will import community created dashboard:

- Click '+' button on left panel and select **'Import'**
- Enter **3131** dashboard id under [Grafana.com](https://grafana.com) Dashboard & click **'Load'**.
- Leave the defaults, select **'Prometheus'** as the endpoint under prometheus data sources drop down, click **'Import'**.

This will show monitoring dashboard for all cluster nodes:



For creating dashboard to monitor all pods, repeat same process as above and enter **3146** for dashboard id:



Source for these instructions - AWS EKS workshop

https://eksworkshop.com/intermediate/240_monitoring/